Revision of week 3 and 4

Nhung Nguyen and Viktor Schlegel

Week 3 - Distributional semantics

We shall know a word by the company it keeps.

Firth (1957)

Week 3 - Count-based approach

- Term-document matrix

tf-idf

	As You Like It	Twelfth Night	Julius Caesar	Henry V
battle	1	0	7	13
good	114	80	62	89
fool	36	58	1	4
wit	20	15	2	3

- Term-term matrix

smoothing

	aardvark	computer	data	pinch	result	sugar	…
apricot	0	0	0	1	0	1	
pineapple	0	0	0	1	0	1	
digital	0	2	1	0	1	0	
information	0	1	6	0	4	0	

Week 3 - Prediction-based approach

- Bengio's language model

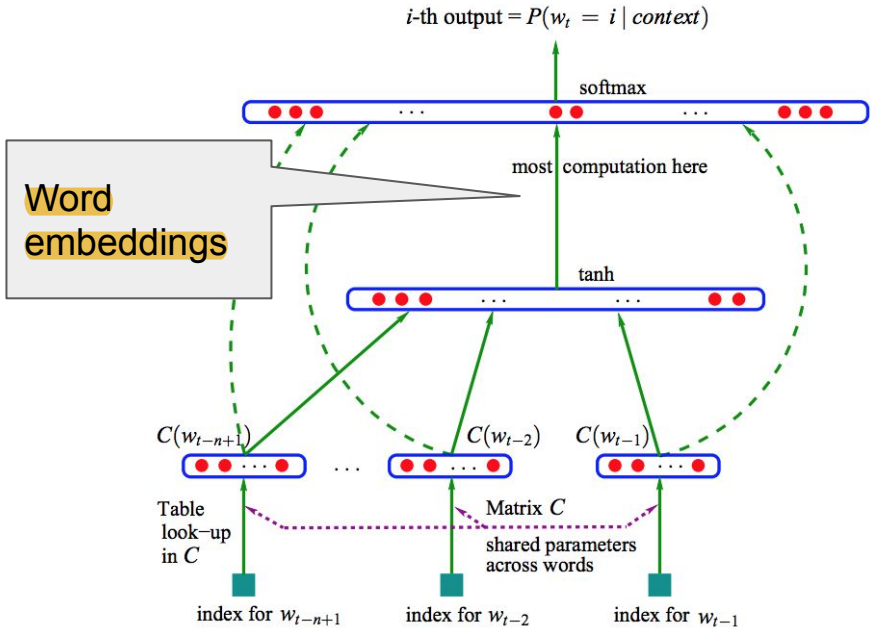

$Pr(w_t|w_{t-1},...,w_{t-m+1}) = \text{softmax}(\mathbf{Wy})$

$\mathbf{y} = \tanh(\mathbf{Vx})$

$\mathbf{x} = \text{concat}(\mathbf{w}_{t-1},...,\mathbf{w}_{t-m+1})$

上下文的

Each *contextual* word w_{t-j} is represented a column of matrix **C**

Week 3 - Prediction-based approach (Cont.)

- Mikolov's CBOW

$\Pr(w_t|w_{t-1},\ldots,w_{t-m+1}) = \text{softmax}(\mathbf{Wy})$

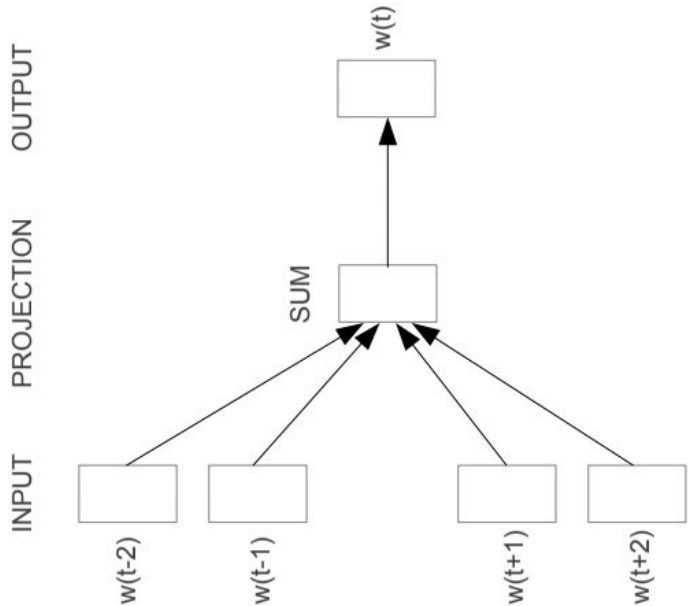

- Mikolov's Skip-gram

$\Pr(.|w_t) = \text{softmax}(\mathbf{Wy})$

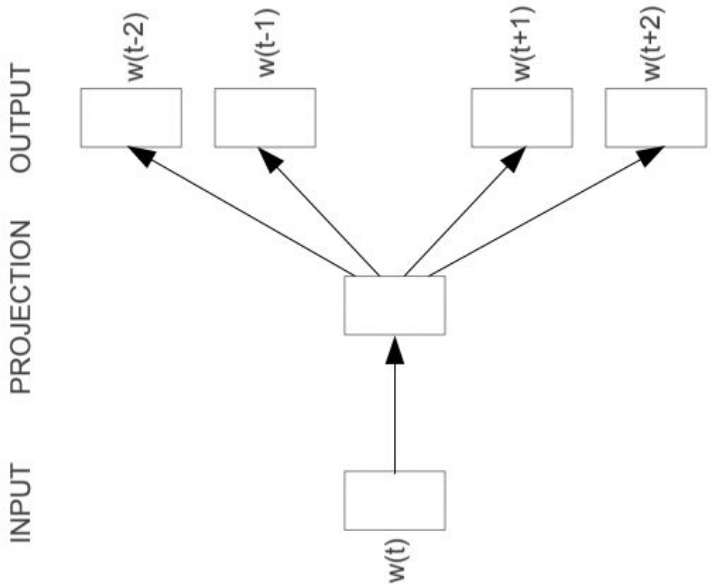

Week 3 - Word Sense Disambiguation

- WordNet: a database of lexical relations
  - A word has different senses

mouse^{1} : .... a *mouse* controlling a computer system in 1968.
mouse^{2} : .... a quiet animal like a *mouse*
bank^{1} : ...a *bank* can hold the investments in a custodial account ...
bank^{2} : ...as agriculture burgeons on the east *bank*, the river ...

- Disambiguating word sense using Lesk's algorithm and supervised learning

Week 4 - Sequence modelling

Sequence classification

S1: I like trains.
S2: The train is arriving on time.

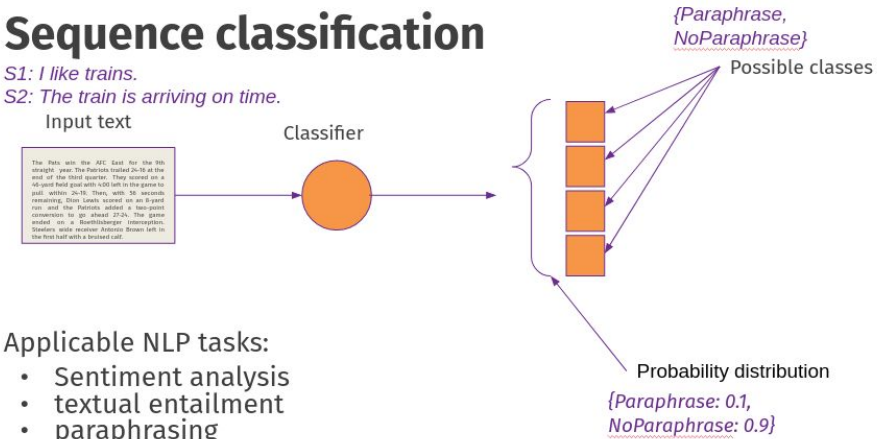

Applicable NLP tasks:

- Sentiment analysis
- textual entailment
- paraphrasing
- question type classification

Span extraction

P: I like trains.
Q: Who likes trains?

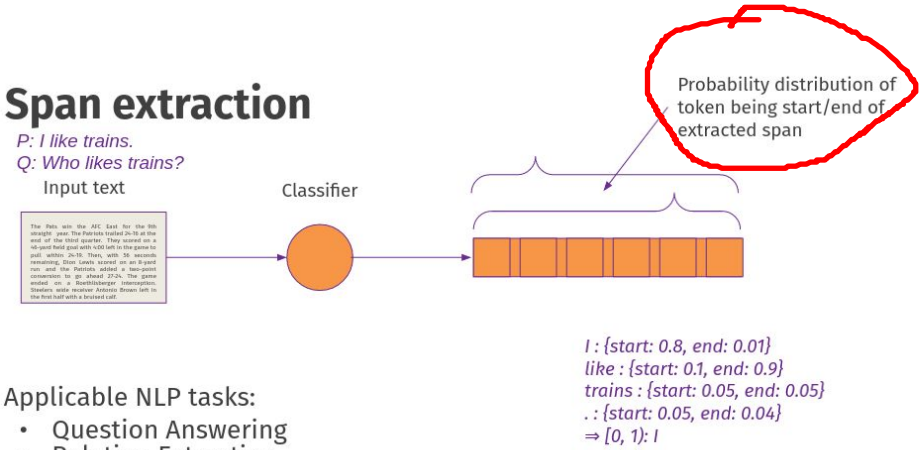

Applicable NLP tasks:

- Question Answering
- Relation Extraction

Week 4 - Sequence modelling

Sequence to sequence

I like trains.

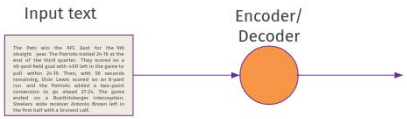

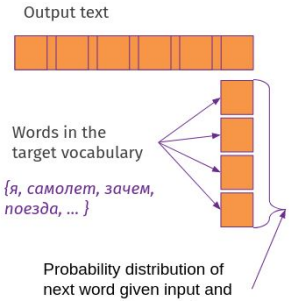

я : {я: 0.85, ты: 0.01, мы: ...
люблю : {люблю: 0.7, пошел: .... }
поезда : {поезда: 0.6, тебя: 0.2 ... }
. : { . : 0.99, UH: 0.001 ... }
⇒ я люблю поезда.

Applicable NLP tasks:

- Translation
- Abstractive Summarisation
- Text generation
- Question answering

Sequence labelling

I like trains.

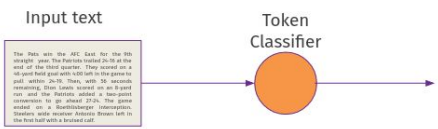

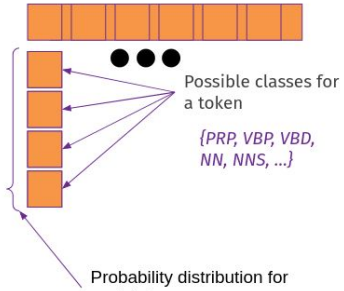

I : {PRP: 0.85, NN: 0.01, DET: 0.01 ...
like : {VBP: 0.7, VBD: 0.2, ... }
trains : {NNS: 0.6, NN: 0.2 ... }
. : { . : 0.99, UH: 0.001 ... }
⇒ I/PRP like/VBP trains/NNS ./.

Applicable NLP tasks:

- POS tagging
- Named Entity Recognition
- OpenIE, Semantic Role Labelling
- question type classification

Week 4: RNN, LSTM, BiLSTM

RNN: Forward run

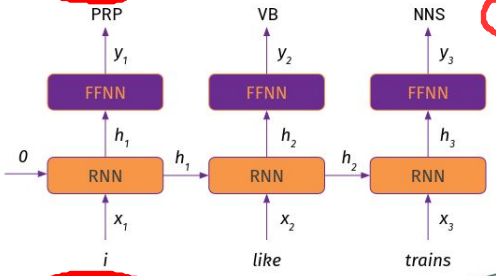

RNN: Backward run

Backpropagation through time

h_t depends on h_{t-1}

Unroll computation graph and do standard backpropagation, because t is not infinite.

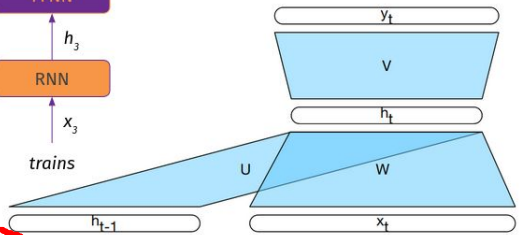

$$h_0 = 0$$
$$h_t = g(Uh_{t-1} + Wx_t + b)$$
$$y_t = f(Vh_t + b_z)$$

RNN: vanishing gradients

< 0.25! Lots of multiplications!

BiRNN

Idea: If we can go from left to right ("past"), why not also just go right to left ("future")

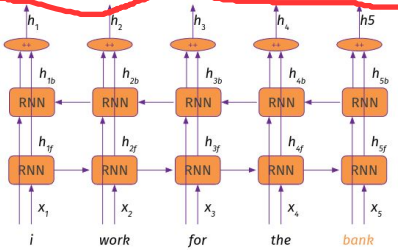

concatenation

$$h_t = [h_{t_f}; h_{t_b}]$$
$$dim(h_t) = dim(h_{t_f}) + dim(h_{t_b})$$

16

LSTM: Context vector

- Context, or memory vector C_t in addition to h_t
- Context information from "the past" calculations
- At any step t, LSTM learns how much of h_t is to be "added" to C_t and how much of C_{t-1} is "kept"

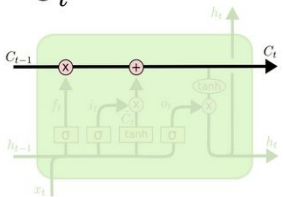

For example:
Information about grammatical gender of subject

Week 4 - Contextualised embeddings

```
d[word] = vector
f(word, context) = contextualised_vector
```

Language modelling with BiLSTM: ELMo

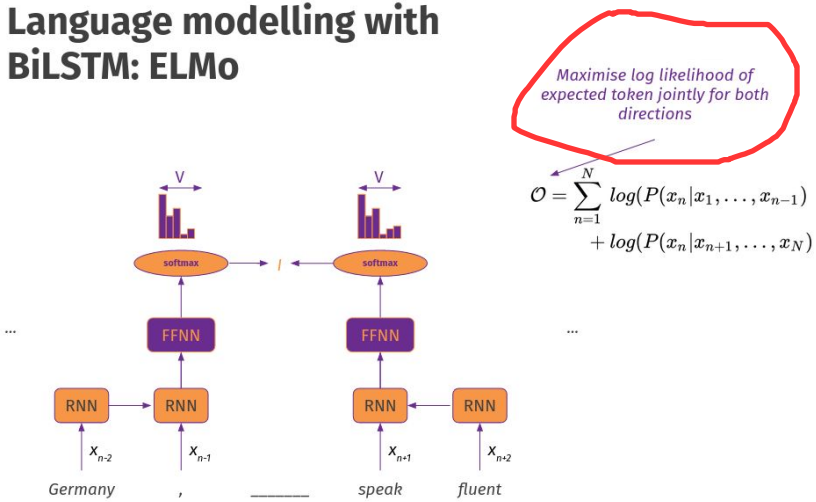

Task: probability of next word given *n* previous words.

$$P(x_{n+1}|x_1, \ldots, x_n)$$

"I grew up in Germany, i speak fluent ____"

syntax information ✔
semantic information ✔
self supervision ✔

Week 4 - NER

NER as a tagging problem (BIO scheme)

	Adam	Smith	works	for	IBM	in	London	.
POS tagging	NNP	NNP	VBZ	IN	NNP	IN	NNP	.
Entity recognition	B_PER	I_PER	O	O	B_ORG	O	B_LOC	O

Begin the mention

Inside the mention

- # classes = 2 * # entity types + 1

Week 4 - NER approaches

- Local approach: tags are ***independent*** each other
  - Any classifiers can be used
    - SVM: truly local
    - RNN, LSTM, BiLSTM: not truly local

- Global approach: tags are ***dependent*** each other
  - Hidden Markov Model (HMM)
  - Conditional Random Fields (CRF)

Week 4: CRF vs. Neural networks

CRF

- Feature engineering
- Do not need pre-trained vectors
- Models are roughly interpretable
- *Perform well with datasets that have many NE categories(*)*

Neural networks

- Do not need features
- Need pre-trained vectors from big language models
- Models use implicit features (created by hidden layers) → not easy to interpret
- *Perform not so well with datasets that have many NE categories(*)*

(*) This observation is only based on my personal experiences

Week 4 - NEL

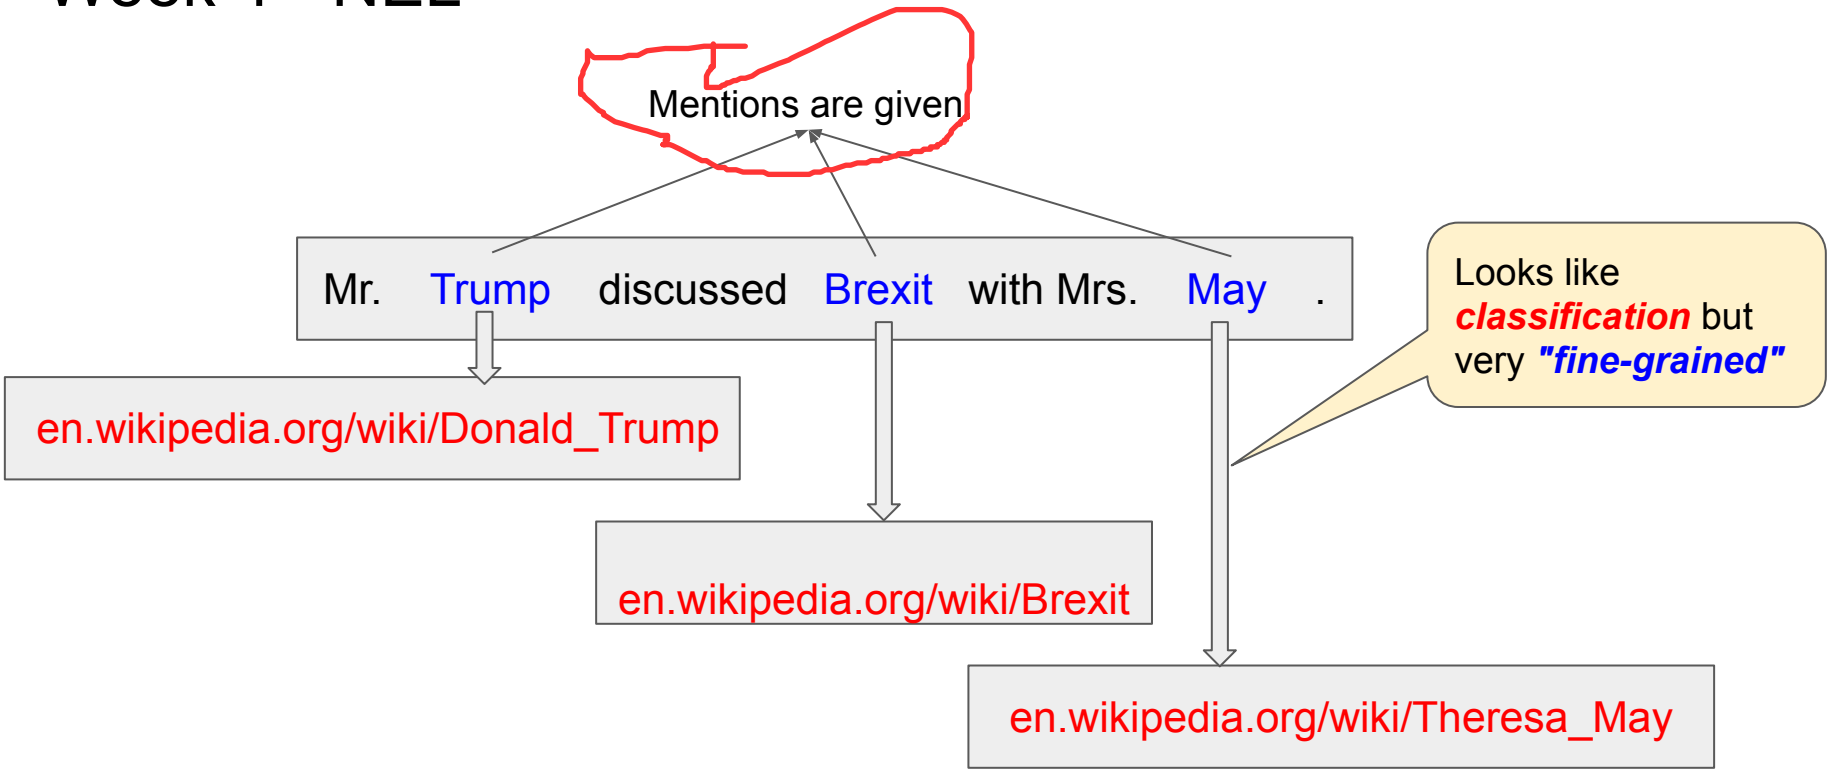
Mentions are given
Mr. Trump discussed Brexit with Mrs. May .
en.wikipedia.org/wiki/Donald_Trump
en.wikipedia.org/wiki/Brexit
en.wikipedia.org/wiki/Theresa_May
Looks like classification but very "fine-grained"

Week 4 - NEL basic steps

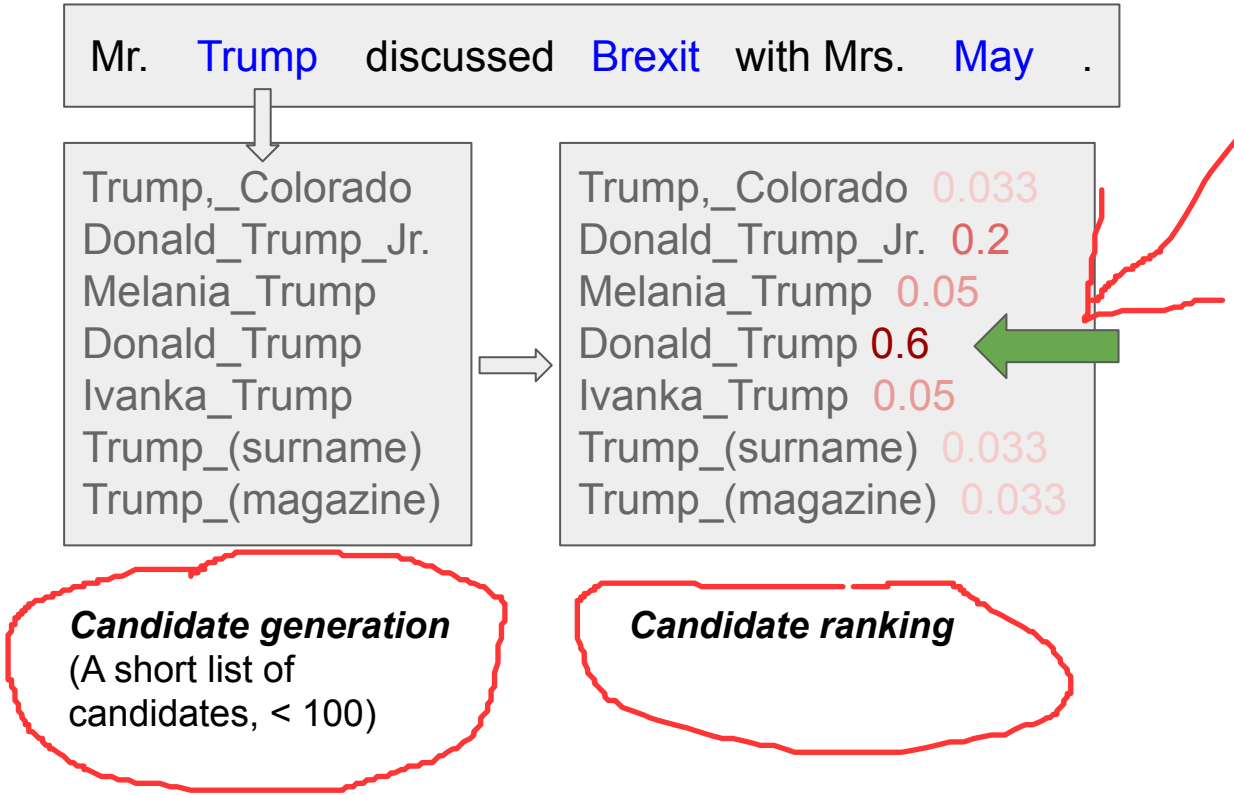

Thank you very much
Good luck with your exams!