

Week 3 (Cont.)
Word Embeddings:
Count-based Approach

Nhung Nguyen
slides courtesy of Phong Le

Recap

- A brief introduction to distributional semantics
- A very simple method using co-occurrence vectors
- Pros and cons of using vectors to represent word meanings

Count-based approach: Term-document matrix

- If word u appears in document d , d is a context of u
1. Collect a lot of documents (from, e.g. Wikipedia)
 2. Count how many time a word u appearing with a document d
 3. The meaning of word u is vector $[\text{count}(u, d_1), \text{count}(u, d_2), \dots]$

	As You Like It	Twelfth Night	Julius Caesar	Henry V
battle	1	0	7	13
good	114	80	62	89
fool	36	58	1	4
wit	20	15	2	3

Shakespeare plays

Term-document matrix (cont.): Document vectors

Two ways to extract information from the matrix

1. Column-wise: a *document* is represented by a $|V|$ -dim vector (V: vocabulary)

	As You Like It	Twelfth Night	Julius Caesar	Henry V
battle	1	0	7	13
good	114	80	62	89
fool	36	58	1	4
wit	20	15	2	3

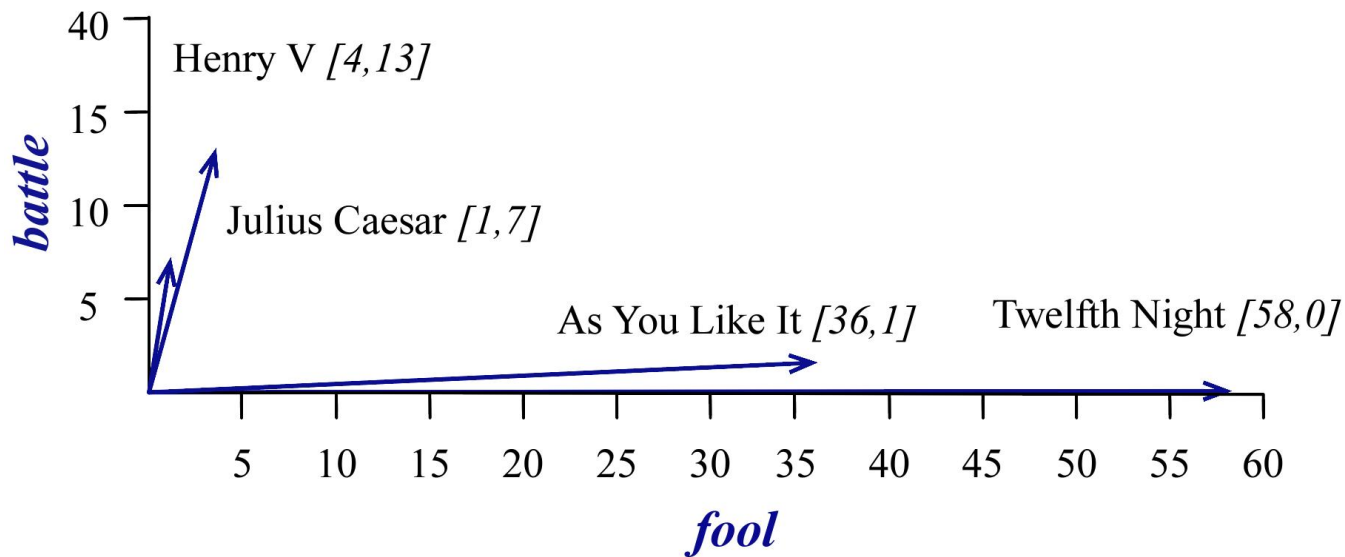
Widely used in information retrieval:

- find similar documents
- find documents close to a query

Term-document matrix (cont.): Document vectors

- Find similar documents:


Two documents that are similar will tend to have similar words



Term-document matrix (cont.): Document vectors

- Find documents close to a query

Consider a query as a document



Term-document matrix (cont.): Word vectors

2. Row-wise: a *word* is represented by a $|D|$ -dim vector (D: document set)

	As You Like It	Twelfth Night	Julius Caesar	Henry V
battle	1	0	7	13
good	114	80	62	89
fool	36	58	1	4
wit	20	15	2	3

Count-based approach: Term-term matrix

- we have seen it before (co-occurrence vectors): Count how many times a word *u* appearing with a word *v*

sugar, a sliced lemon, a tablespoonful of
their enjoyment. Cautiously she sampled her first
well suited to programming on the digital
for the purpose of gathering data and

apricot
pineapple
computer.
information

jam, a pinch each of,
and another fruit whose taste she likened
In finding the optimal R-stage policy from
necessary for the study authorized in the

	aardvark	computer	data	pinch	result	sugar	...
apricot	0	0	0	1	0	1	
pineapple	0	0	0	1	0	1	
digital	0	2	1	0	1	0	
information	0	1	6	0	4	0	

Count-based approach: raw frequency is bad

sugar, a sliced lemon, a tablespoonful of	apricot	jam, a pinch each of,
their enjoyment. Cautiously she sampled her first	pineapple	and another fruit whose taste she likened
well suited to programming on the digital	computer.	In finding the optimal R-stage policy from
for the purpose of gathering data and	information	necessary for the study authorized in the

- Not all contextual words are equally important: of, a, ... vs. sugar, jam, fruit...
- Which words are important, which ones are not?
 - infrequent words are more important than frequent ones (examples?)
 - correlated words are more important than uncorrelated ones (examples?)
 - ...

→ weighing schemes (TF-IDF, PMI,...)

Weighing terms: **TF-IDF** (for term-document matrix)

- tf (term frequency): frequency count

Frequency of term t in document d

$$tf(t, d) = \log_{10}(1 + count(t, d))$$

- idf (inverse document frequency): popular terms (terms that appear in many documents) are down weighed

Total number of documents in the collection

$$idf(t) = \log_{10} \frac{N}{df(t)}$$

Number of documents in which term t occurs

TF-IDF: $tf-idf(t, d) = tf(t, d) \times idf(t)$

Weighing terms: TF-IDF (example)

Shakespeare plays

Word	df	idf
Romeo	1	1.57
salad	2	1.27
Falstaff	4	0.967
forest	12	0.489
battle	21	0.246
wit	34	0.037
fool	36	0.012
good	37	0
sweet	37	0

	As You Like It	Twelfth Night	Julius Caesar	Henry V
battle	1	0	7	13
good	114	80	62	89
fool	36	58	1	4
wit	20	15	2	3

raw frequency

tf-idf

	As You Like It	Twelfth Night	Julius Caesar	Henry V
battle	0.074	0	0.22	0.28
good	0	0	0	0
fool	0.019	0.021	0.0036	0.0083
wit	0.049	0.044	0.018	0.022

Count-based approach: so many 0s

	aardvark	computer	data	pinch	result	sugar	...
apricot	0	0	0	1	0	1	
pineapple	0	0	0	1	0	1	
digital	0	2	1	0	1	0	
information	0	1	6	0	4	0	

- Many word pairs should have > 0 counts, but their corresponding matrix entries are 0s because of lacking data (**data sparsity**)
 - **Laplace smoothing**: adding 1 to every entry (pseudocount)

Pros

- Simple and intuitive
- Dimensions are meaningful (e.g. each dim is a document / a contextual word)
→ easy to debug and interpret (*Think about Explainable AI*)

Cons

- Word/document vectors are **sparse** (dims are $|V|$, vocabulary size, or $|D|$, number of documents, often from 2k to 10k) → **difficult for machine learning algorithms**
- How to represent word meaning in a specific context?

From sparse vectors to dense vectors

- Employ dimensionality reduction (e.g. latent semantic analysis - LSA)
- Use a different approach: prediction (coming up next)